

✓ NER task - Product Attribute Extraction

```
!pip install unsloth
```



Collecting unsloth

Downloading unsloth-2025.7.3-py3-none-any.whl.metadata (47 kB)

47.2/47.2 kB 2.8 MB/s eta 0:00:00

Collecting unsloth_zoo>=2025.7.4 (from unsloth)

Downloading unsloth_zoo-2025.7.4-py3-none-any.whl.metadata (8.1 kB)

Requirement already satisfied: torch>=2.4.0 in /usr/local/lib/python3.11/dist-packages (from unsloth) (2.6.0+cu124)

Collecting xformers>=0.0.27.post2 (from unsloth)

Downloading xformers-0.0.31.post1-cp39-abi3-manylinux_2_28_x86_64.whl.metadata (1.1 kB)

Collecting bitsandbytes (from unsloth)

Downloading bitsandbytes-0.46.1-py3-none-manylinux_2_24_x86_64.whl.metadata (10 kB)

Requirement already satisfied: triton>=3.0.0 in /usr/local/lib/python3.11/dist-packages (from unsloth) (3.2.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from unsloth) (24.2)

Collecting tyro (from unsloth)

Downloading tyro-0.9.26-py3-none-any.whl.metadata (12 kB)

Requirement already satisfied: transformers!=4.47.0,!=4.52.0,!=4.52.1,!=4.52.2,!=4.52.3,!=4.53.0,>=4.51.3 in /usr/local/lib/python3.11/dist-packages (from unsloth) (4.51.3)

Collecting datasets<4.0.0,>=3.4.1 (from unsloth)

Downloading datasets-3.6.0-py3-none-any.whl.metadata (19 kB)

Requirement already satisfied: sentencepiece>=0.2.0 in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.2.0)

Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from unsloth) (4.67.1)

Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packages (from unsloth) (5.9.5)

Requirement already satisfied: wheel>=0.42.0 in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.45.1)

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from unsloth) (2.0.2)

Requirement already satisfied: accelerate>=0.34.1 in /usr/local/lib/python3.11/dist-packages (from unsloth) (1.8.1)

Collecting trl!=0.15.0,!=0.9.0,!=0.9.1,!=0.9.2,!=0.9.3,>=0.7.9 (from unsloth)

Downloading trl-0.19.1-py3-none-any.whl.metadata (10 kB)

Requirement already satisfied: peft!=0.11.0,>=0.7.1 in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.16.0)

Requirement already satisfied: protobuf in /usr/local/lib/python3.11/dist-packages (from unsloth) (5.29.5)

Requirement already satisfied: huggingface_hub in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.33.2)

Requirement already satisfied: hf_transfer in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.1.9)

Requirement already satisfied: diffusers in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.34.0)

Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (from unsloth) (0.21.0+cu124)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from accelerate>=0.34.1->unsloth) (6.0.2)

Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.11/dist-packages (from accelerate>=0.34.1->unsloth) (0.5.3)

Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (3.18.0)

Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (18.1.0)

Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (0.3.7)

Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (2.2.2)

Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (2.32.3)

Requirement already satisfied: xxhash in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (3.5.0)

Requirement already satisfied: multiprocessing<0.70.17 in /usr/local/lib/python3.11/dist-packages (from datasets<4.0.0,>=3.4.1->unsloth) (0.70.15)

Collecting fsspec<=2025.3.0,>=2023.1.0 (from fsspec[http]<=2025.3.0,>=2023.1.0->datasets<4.0.0,>=3.4.1->unsloth)

Downloading fsspec-2025.3.0-py3-none-any.whl.metadata (11 kB)

Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->unsloth) (4.14.1)

Requirement already satisfied: hf_xet<2.0.0,>=1.1.2 in /usr/local/lib/python3.11/dist-packages (from huggingface_hub->unsloth) (1.1.5)

Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch>=2.4.0->unsloth) (3.5)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch>=2.4.0->unsloth) (3.1.6)

Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch>=2.4.0->unsloth)

Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)

Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch>=2.4.0->unsloth)

Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch>=2.4.0->unsloth)
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch>=2.4.0->unsloth)
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch>=2.4.0->unsloth)
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch>=2.4.0->unsloth)
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch>=2.4.0->unsloth)
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch>=2.4.0->unsloth)
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cuspars-cu12==12.3.1.170 (from torch>=2.4.0->unsloth)
Downloading nvidia_cuspars-cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch>=2.4.0->unsloth) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch>=2.4.0->unsloth) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.4.0->unsloth) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch>=2.4.0->unsloth)
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch>=2.4.0->unsloth) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch>=2.4.0->unsloth) (1.3.0)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from transformers!=4.47.0,!=4.52.0,!=4.52.1,!=4.52.2,
Requirement already satisfied: tokenizers<0.22,>=0.21 in /usr/local/lib/python3.11/dist-packages (from transformers!=4.47.0,!=4.52.0,!=4.52.1,!=4.52.2,
Collecting protobuf (from unsloth)
Downloading protobuf-3.20.3-py2.py3-none-any.whl.metadata (720 bytes)
Collecting cut_cross_entropy (from unsloth_zoo>=2025.7.4->unsloth)
Downloading cut_cross_entropy-25.1.1-py3-none-any.whl.metadata (9.3 kB)
Requirement already satisfied: pillow in /usr/local/lib/python3.11/dist-packages (from unsloth_zoo>=2025.7.4->unsloth) (11.2.1)
Collecting msgspec (from unsloth_zoo>=2025.7.4->unsloth)
Downloading msgspec-0.19.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (6.9 kB)
Collecting torch>=2.4.0 (from unsloth)
Downloading torch-2.7.1-cp311-cp311-manylinux_2_28_x86_64.whl.metadata (29 kB)
Collecting sympy>=1.13.3 (from torch>=2.4.0->unsloth)
Downloading sympy-1.14.0-py3-none-any.whl.metadata (12 kB)
Collecting nvidia-cuda-nvrtc-cu12==12.6.77 (from torch>=2.4.0->unsloth)
Downloading nvidia_cuda_nvrtc_cu12-12.6.77-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.6.77 (from torch>=2.4.0->unsloth)
Downloading nvidia_cuda_runtime_cu12-12.6.77-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.6.80 (from torch>=2.4.0->unsloth)
Downloading nvidia_cuda_cupti_cu12-12.6.80-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.5.1.17 (from torch>=2.4.0->unsloth)
Downloading nvidia_cudnn_cu12-9.5.1.17-py3-none-manylinux_2_28_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12 (from nvidia-cudnn-cu12==9.1.0.70->torch>=2.4.0->unsloth)
Downloading nvidia_cublas_cu12-12.6.4.1-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.3.0.4 (from torch>=2.4.0->unsloth)
Downloading nvidia_cufft_cu12-11.3.0.4-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.7.77 (from torch>=2.4.0->unsloth)
Downloading nvidia_curand_cu12-10.3.7.77-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl.metadata (1.5 kB)

```
Collecting nvidia-cusolver-cu12==11.7.1.2 (from torch>=2.4.0->unsloth)
  Downloading nvidia_cusolver_cu12-11.7.1.2-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cuspars-cu12 (from nvidia-cusolver-cu12==11.6.1.9->torch>=2.4.0->unsloth)
  Downloading nvidia_cuspars-cu12-12.5.4.2-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==0.6.3 (from torch>=2.4.0->unsloth)
  Downloading nvidia_cusparselt_cu12-0.6.3-py3-none-manylinux2014_x86_64.whl.metadata (6.8 kB)
Collecting nvidia-nccl-cu12==2.26.2 (from torch>=2.4.0->unsloth)
  Downloading nvidia_nccl_cu12-2.26.2-py3-none-manylinux2014_x86_64.whl.metadata (2.0 kB)
Collecting nvidia-nvtx-cu12==12.6.77 (from torch>=2.4.0->unsloth)
  Downloading nvidia_nvtx_cu12-12.6.77-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-nvjitlink-cu12 (from nvidia-cusolver-cu12==11.6.1.9->torch>=2.4.0->unsloth)
  Downloading nvidia_nvjitlink_cu12-12.6.85-py3-none-manylinux2010_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufile-cu12==1.11.1.6 (from torch>=2.4.0->unsloth)
  Downloading nvidia_cufile_cu12-1.11.1.6-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting triton>=3.0.0 (from unsloth)
  Downloading triton-3.3.1-cp311-cp311-manylinux_2_27_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: setuptools>=40.8.0 in /usr/local/lib/python3.11/dist-packages (from triton>=3.0.0->unsloth) (75.2.0)
Requirement already satisfied: importlib_metadata in /usr/local/lib/python3.11/dist-packages (from diffusers->unsloth) (8.7.0)
INFO: pip is looking at multiple versions of torchvision to determine which version is compatible with other requirements. This could take a while
Collecting torchvision (from unsloth)
  Downloading torchvision-0.22.1-cp311-cp311-manylinux_2_28_x86_64.whl.metadata (6.1 kB)
Requirement already satisfied: docstring-parser>=0.15 in /usr/local/lib/python3.11/dist-packages (from tyro->unsloth) (0.16)
Requirement already satisfied: rich>=11.1.0 in /usr/local/lib/python3.11/dist-packages (from tyro->unsloth) (13.9.4)
Collecting shtab>=1.5.6 (from tyro->unsloth)
  Downloading shtab-1.7.2-py3-none-any.whl.metadata (7.4 kB)
Requirement already satisfied: typeguard>=4.0.0 in /usr/local/lib/python3.11/dist-packages (from tyro->unsloth) (4.4.4)
Requirement already satisfied: aiohttp!=4.0.0a0,!4.0.0a1 in /usr/local/lib/python3.11/dist-packages (from fsspec[http]<=2025.3.0,>=2023.1.0->data)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests>=2.32.2->datasets<4.0.0,>=3.4.1-)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests>=2.32.2->datasets<4.0.0,>=3.4.1->unsloth) (3)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests>=2.32.2->datasets<4.0.0,>=3.4.1->unsloth) (3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests>=2.32.2->datasets<4.0.0,>=3.4.1->unsloth) (2025.2)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich>=11.1.0->tyro->unsloth) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich>=11.1.0->tyro->unsloth) (2.19.2)
Requirement already satisfied: zipp>=3.20 in /usr/local/lib/python3.11/dist-packages (from importlib_metadata->diffusers->unsloth) (3.23.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2->torch>=2.4.0->unsloth) (3.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas->datasets<4.0.0,>=3.4.1->unsloth) (2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas->datasets<4.0.0,>=3.4.1->unsloth) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas->datasets<4.0.0,>=3.4.1->unsloth) (2025.2)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.11/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]<=2025.3)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich>=11.1.0->tyro->unsloth) (0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas->datasets<4.0.0,>=3.4.1->u)
Downloading unsloth-2025.7.3-py3-none-any.whl (297 kB)
```

297.5/297.5 kB 12.2 MB/s eta 0:00:00

```
Downloading datasets-3.6.0-py3-none-any.whl (491 kB)
_____ 491.5/491.5 kB 26.8 MB/s eta 0:00:00
Downloading trl-0.19.1-py3-none-any.whl (376 kB)
_____ 376.2/376.2 kB 34.3 MB/s eta 0:00:00
Downloading unsloth_zoo-2025.7.4-py3-none-any.whl (162 kB)
_____ 162.8/162.8 kB 16.8 MB/s eta 0:00:00
Downloading protobuf-3.20.3-py2.py3-none-any.whl (162 kB)
_____ 162.1/162.1 kB 16.5 MB/s eta 0:00:00
Downloading xformers-0.0.31.post1-cp39-abi3-manylinux_2_28_x86_64.whl (117.1 MB)
_____ 117.1/117.1 MB 10.0 MB/s eta 0:00:00
Downloading torch-2.7.1-cp311-cp311-manylinux_2_28_x86_64.whl (821.2 MB)
_____ 821.2/821.2 MB 1.3 MB/s eta 0:00:00
Downloading triton-3.3.1-cp311-cp311-manylinux_2_27_x86_64.manylinux_2_28_x86_64.whl (155.7 MB)
_____ 155.7/155.7 MB 7.9 MB/s eta 0:00:00
Downloading nvidia_cublas_cu12-12.6.4.1-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (393.1 MB)
_____ 393.1/393.1 MB 5.2 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.6.80-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (8.9 MB)
_____ 8.9/8.9 MB 109.8 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.6.77-py3-none-manylinux2014_x86_64.whl (23.7 MB)
_____ 23.7/23.7 MB 93.1 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.6.77-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (897 kB)
_____ 897.7/897.7 kB 59.7 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.5.1.17-py3-none-manylinux_2_28_x86_64.whl (571.0 MB)
_____ 571.0/571.0 MB 3.2 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.3.0.4-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (200.2 MB)
_____ 200.2/200.2 MB 7.3 MB/s eta 0:00:00
Downloading nvidia_cufile_cu12-1.11.1.6-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (1.1 MB)
_____ 1.1/1.1 MB 64.1 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.7.77-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (56.3 MB)
_____ 56.3/56.3 MB 14.7 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.7.1.2-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (158.2 MB)
_____ 158.2/158.2 MB 7.8 MB/s eta 0:00:00
Downloading nvidia_cusparsesearch_cu12-12.5.4.2-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (216.6 MB)
_____ 216.6/216.6 MB 6.1 MB/s eta 0:00:00
Downloading nvidia_cusparselt_cu12-0.6.3-py3-none-manylinux2014_x86_64.whl (156.8 MB)
_____ 156.8/156.8 MB 7.7 MB/s eta 0:00:00
Downloading nvidia_nccl_cu12-2.26.2-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (201.3 MB)
_____ 201.3/201.3 MB 1.6 MB/s eta 0:00:00
Downloading nvidia_nvjitlink_cu12-12.6.85-py3-none-manylinux2010_x86_64.manylinux_2_12_x86_64.whl (19.7 MB)
_____ 19.7/19.7 MB 96.4 MB/s eta 0:00:00
Downloading nvidia_nvtx_cu12-12.6.77-py3-none-manylinux2014_x86_64.manylinux_2_17_x86_64.whl (89 kB)
_____ 89.3/89.3 kB 8.7 MB/s eta 0:00:00
Downloading bitsandbytes-0.46.1-py3-none-manylinux_2_24_x86_64.whl (72.9 MB)
_____ 72.9/72.9 MB 12.0 MB/s eta 0:00:00
Downloading torchvision-0.22.1-cp311-cp311-manylinux_2_28_x86_64.whl (7.5 MB)
_____ 7.5/7.5 MB 112.0 MB/s eta 0:00:00
Downloading tyro-0.9.26-py3-none-any.whl (128 kB)
_____ 129.0/129.0 kB 12.8 MB/s eta 0:00:00
Downloading fsspec-2025.3.0-py3-none-any.whl (193 kB)
_____ 193.6/193.6 kB 19.0 MB/s eta 0:00:00
```

```
193.6/193.6 kB 19.0 MB/s eta 0:00:00
Downloading shtab-1.7.2-py3-none-any.whl (14 kB)
Downloading sympy-1.14.0-py3-none-any.whl (6.3 MB)
6.3/6.3 MB 106.1 MB/s eta 0:00:00
Downloading cut_cross_entropy-25.1.1-py3-none-any.whl (22 kB)
Downloading msgspec-0.19.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (210 kB)
210.7/210.7 kB 21.3 MB/s eta 0:00:00
Installing collected packages: nvidia-cusparselt-cu12, triton, sympy, shtab, protobuf, nvidia-nvtx-cu12, nvidia-nvjitlink-cu12, nvidia-nccl-cu12,
Attempting uninstall: nvidia-cusparselt-cu12
  Found existing installation: nvidia-cusparselt-cu12 0.6.2
  Uninstalling nvidia-cusparselt-cu12-0.6.2:
    Successfully uninstalled nvidia-cusparselt-cu12-0.6.2
Attempting uninstall: triton
  Found existing installation: triton 3.2.0
  Uninstalling triton-3.2.0:
    Successfully uninstalled triton-3.2.0
Attempting uninstall: sympy
  Found existing installation: sympy 1.13.1
  Uninstalling sympy-1.13.1:
    Successfully uninstalled sympy-1.13.1
Attempting uninstall: protobuf
  Found existing installation: protobuf 5.29.5
  Uninstalling protobuf-5.29.5:
    Successfully uninstalled protobuf-5.29.5
Attempting uninstall: nvidia-nvtx-cu12
  Found existing installation: nvidia-nvtx-cu12 12.4.127
  Uninstalling nvidia-nvtx-cu12-12.4.127:
    Successfully uninstalled nvidia-nvtx-cu12-12.4.127
Attempting uninstall: nvidia-nvjitlink-cu12
  Found existing installation: nvidia-nvjitlink-cu12 12.5.82
  Uninstalling nvidia-nvjitlink-cu12-12.5.82:
    Successfully uninstalled nvidia-nvjitlink-cu12-12.5.82
Attempting uninstall: nvidia-nccl-cu12
  Found existing installation: nvidia-nccl-cu12 2.21.5
  Uninstalling nvidia-nccl-cu12-2.21.5:
    Successfully uninstalled nvidia-nccl-cu12-2.21.5
Attempting uninstall: nvidia-curand-cu12
  Found existing installation: nvidia-curand-cu12 10.3.6.82
  Uninstalling nvidia-curand-cu12-10.3.6.82:
    Successfully uninstalled nvidia-curand-cu12-10.3.6.82
Attempting uninstall: nvidia-cuda-runtime-cu12
  Found existing installation: nvidia-cuda-runtime-cu12 12.5.82
  Uninstalling nvidia-cuda-runtime-cu12-12.5.82:
    Successfully uninstalled nvidia-cuda-runtime-cu12-12.5.82
Attempting uninstall: nvidia-cuda-nvrtc-cu12
  Found existing installation: nvidia-cuda-nvrtc-cu12 12.5.82
  Uninstalling nvidia-cuda-nvrtc-cu12-12.5.82:
    Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.5.82
Attempting uninstall: nvidia-cuda-cupti-cu12
  Found existing installation: nvidia-cuda-cupti-cu12 12.5.82
```



```
Found existing installation: nvidia-cuda-cupti-cu12-12.5.82
Uninstalling nvidia-cuda-cupti-cu12-12.5.82:
  Successfully uninstalled nvidia-cuda-cupti-cu12-12.5.82
Attempting uninstall: nvidia-cublas-cu12
  Found existing installation: nvidia-cublas-cu12 12.5.3.2
  Uninstalling nvidia-cublas-cu12-12.5.3.2:
    Successfully uninstalled nvidia-cublas-cu12-12.5.3.2
Attempting uninstall: fsspec
  Found existing installation: fsspec 2025.3.2
  Uninstalling fsspec-2025.3.2:
    Successfully uninstalled fsspec-2025.3.2
Attempting uninstall: nvidia-cuspars-cu12
  Found existing installation: nvidia-cuspars-cu12 12.5.1.3
  Uninstalling nvidia-cuspars-cu12-12.5.1.3:
    Successfully uninstalled nvidia-cuspars-cu12-12.5.1.3
Attempting uninstall: nvidia-cufft-cu12
  Found existing installation: nvidia-cufft-cu12 11.2.3.61
  Uninstalling nvidia-cufft-cu12-11.2.3.61:
    Successfully uninstalled nvidia-cufft-cu12-11.2.3.61
Attempting uninstall: nvidia-cudnn-cu12
  Found existing installation: nvidia-cudnn-cu12 9.3.0.75
  Uninstalling nvidia-cudnn-cu12-9.3.0.75:
    Successfully uninstalled nvidia-cudnn-cu12-9.3.0.75
Attempting uninstall: nvidia-cusolver-cu12
  Found existing installation: nvidia-cusolver-cu12 11.6.3.83
  Uninstalling nvidia-cusolver-cu12-11.6.3.83:
    Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Attempting uninstall: torch
  Found existing installation: torch 2.6.0+cu124
  Uninstalling torch-2.6.0+cu124:
    Successfully uninstalled torch-2.6.0+cu124
Attempting uninstall: datasets
  Found existing installation: datasets 2.14.4
  Uninstalling datasets-2.14.4:
    Successfully uninstalled datasets-2.14.4
Attempting uninstall: torchvision
  Found existing installation: torchvision 0.21.0+cu124
  Uninstalling torchvision-0.21.0+cu124:
    Successfully uninstalled torchvision-0.21.0+cu124
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the fol
fastai 2.7.19 requires torch<2.7,>=1.10, but you have torch 2.7.1 which is incompatible.
grpcio-status 1.71.2 requires protobuf<6.0dev,>=5.26.1, but you have protobuf 3.20.3 which is incompatible.
torchaudio 2.6.0+cu124 requires torch==2.6.0, but you have torch 2.7.1 which is incompatible.
ydf 0.12.0 requires protobuf<6.0.0,>=5.29.1, but you have protobuf 3.20.3 which is incompatible.
tensorflow-metadata 1.17.2 requires protobuf>=4.25.2; python_version >= "3.11", but you have protobuf 3.20.3 which is incompatible.
gcsfs 2025.3.2 requires fsspec==2025.3.2, but you have fsspec 2025.3.0 which is incompatible.
Successfully installed bitsandbytes-0.46.1 cut_cross_entropy-25.1.1 datasets-3.6.0 fsspec-2025.3.0 msgspec-0.19.0 nvidia-cublas-cu12-12.6.4.1 nvid
WARNING: The following packages were previously imported in this runtime:
```

[google,torch,torchgen]

You must restart the runtime in order to use newly installed versions.

RESTART SESSION




```
import os
import json
import torch
from unsloth import FastLanguageModel
from transformers import TrainingArguments, Trainer, DataCollatorForLanguageModeling
from datasets import load_dataset
```

```
print("torch version", torch.__version__)
print("cuda available", torch.cuda.is_available())
```

```
🔄 🦊 Unsloth: Will patch your computer to enable 2x faster free finetuning.
🦊 Unsloth Zoo will now patch everything to make training faster!
torch version 2.7.1+cu126
cuda available True
```

✓ =====

✅ Step 1: Prepare NER-style Dataset

=====

```
dataset = [
    {
        "Instruction": "Extract product attributes from the description",
        "Input": "This matte black case is designed for the iPhone 13 Pro Max. It's made from TPU and polycarbonate, weighs 1.2 ounces, and was manufactured in China.",
        "Output": json.dumps({
            "Compatible Phone Models": "iPhone 13 Pro Max",
            "Color": "matte black",
            "Material": "TPU and polycarbonate",
            "Item Weight": "1.2 ounces",
            "Country of Origin": "China"
        }, indent=2)
    },
    {
        "Instruction": "Extract product attributes from the description",
        "Input": "Made for Samsung Galaxy S22 Ultra, this case comes in sky blue and features a vegan leather finish. It weighs 1.5 ounces and is made in South Korea.",
        "Output": json.dumps({
            "Compatible Phone Models": "Samsung Galaxy S22 Ultra",
            "Color": "sky blue",
            "Material": "vegan leather",
            "Item Weight": "1.5 ounces",
            "Country of Origin": "South Korea"
        }, indent=2)
    }
]
```

```

        "Material": "vegan leather",
        "Item Weight": "1.5 ounces",
        "Country of Origin": "South Korea"
    }, indent=2)
},
{
    "Instruction": "Extract product attributes from the description",
    "Input": "A protective screen cover for the iPad Air 5th Gen, built with 9H tempered glass, this 2.1-ounce product is manufactured in Japan.",
    "Output": json.dumps({
        "Compatible Phone Models": "iPad Air 5th Gen",
        "Material": "9H tempered glass",
        "Item Weight": "2.1 ounces",
        "Country of Origin": "Japan"
    }, indent=2)
},
]

```

Generate synthetic data

```

colors = ["red", "black", "white", "green", "navy blue", "champagne gold"]
models = ["iPhone 14", "Pixel 8 Pro", "OnePlus 11", "Samsung Galaxy A54", "iPad Mini 6"]
materials = ["silicone", "plastic", "TPU", "carbon fiber", "tempered glass", "leather"]
weights = ["1.0 ounces", "1.5 ounces", "2.0 ounces", "2.5 ounces"]
countries = ["China", "India", "Germany", "USA", "Vietnam", "South Korea"]

```

import random

```

for _ in range(27):
    phone = random.choice(models)
    color = random.choice(colors)
    material = random.choice(materials)
    weight = random.choice(weights)
    country = random.choice(countries)
    description = f"This {color} case is compatible with the {phone}, made from {material}. It weighs {weight} and is manufactured in {country}."
    attributes = {
        "Compatible Phone Models": phone,
        "Color": color,
        "Material": material,
        "Item Weight": weight,
        "Country of Origin": country
    }
    dataset.append({
        "Instruction": "Extract product attributes from the description",
        "Input": description,
        "Output": json.dumps(attributes, indent=2)
    })

```

```

}))

os.makedirs("data", exist_ok=True)
with open("data/ner_data.json", "w") as f:
    for item in dataset:
        json_record = json.dumps(item)
        f.write(json_record + "\n")

```

```

print("✅ NER-style sample data saved.")

```

```

🔄 ✅ NER-style sample data saved.

```

▼ =====

✅ Step 2: Load and Prepare Model

=====

```

model_name = "mistralai/Mistral-7B-Instruct-v0.2"

model, tokenizer = FastLanguageModel.from_pretrained(
    model_name=model_name,
    load_in_4bit=True,
)

model = FastLanguageModel.get_peft_model(
    model,
    r=8,
    lora_alpha=16,
    lora_dropout=0.05,
    bias="none",
    target_modules=["q_proj", "v_proj", "k_proj"],
    use_gradient_checkpointing=True,
)

```

```

🔄 ==((====))== Unsloth 2025.7.3: Fast Mistral patching. Transformers: 4.53.1.
    \ \   /|   Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform: Linux.
0^0/ \_/ \   Torch: 2.7.1+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton: 3.3.1

```

```
\      /      Bfloat16 = False. FA [Xformers = 0.0.31.post1. FA2 = False]
"-_____"      Free license: http://github.com/unslothai/unsloth
Unsloth: Fast downloading is enabled - ignore downloading bars which are red colored!
Unsloth: Dropout = 0 is supported for fast patching. You are using dropout = 0.05.
Unsloth will patch all other layers, except LoRA matrices, causing a performance hit.
Unsloth 2025.7.3 patched 32 layers with 0 QKV layers, 0 0 layers and 0 MLP layers.
```

▼ =====

✔ Step 3: Load Dataset & Tokenize

=====

```
dataset = load_dataset("json", data_files="data/ner_data.json", split="train")
```

```
def tokenize_fn(examples):
    prompt = (
        f"### Instruction:\n{examples['Instruction']}\n\n"
        f"### Input:\n{examples['Input']}\n\n"
        f"### Output (in JSON format):\n{examples['Output']}"
    )
    tokenized = tokenizer(
        prompt,
        truncation=True,
        max_length=512,
        padding="max_length",
    )
    tokenized["labels"] = tokenized["input_ids"].copy()
    return tokenized
```

```
tokenized_dataset = dataset.map(tokenize_fn)
```

```
🔄 Generating train split:      30/0 [00:00<00:00, 549.94 examples/s]

Map: 100%                      30/30 [00:00<00:00, 302.75 examples/s]
```

▼ =====

✔ Step 4: Train

=====

```
training_args = TrainingArguments(
    output_dir="finetuned_model",
    per_device_train_batch_size=2,
    gradient_accumulation_steps=2,
    learning_rate=2e-4,
    logging_steps=1,
    num_train_epochs=3,
    optim="adamw_torch",
    lr_scheduler_type="cosine",
    report_to="none"
)

data_collator = DataCollatorForLanguageModeling(
    tokenizer=tokenizer,
    mlm=False,
)

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_dataset,
    data_collator=data_collator
)

trainer.train()
print("✔ Training complete!")

model.save_pretrained("finetuned_model")
tokenizer.save_pretrained("finetuned_model")
print("✔ Model adapters and tokenizer saved.")
```

⇒ ==((====))== Unsloth - 2x faster free finetuning | Num GPUs used = 1
 \ \ / | Num examples = 30 | Num Epochs = 3 | Total steps = 24
0^0/ _/ \ Batch size per device = 2 | Gradient accumulation steps = 2
 \ _____ / Data Parallel GPUs = 1 | Total batch size (2 x 2 x 1) = 4
"-_____" Trainable parameters = 4,718,592 of 7,246,450,688 (0.07% trained)
Unsloth: Will smartly offload gradients to save VRAM!
[24/24 02:07, Epoch 3/3]

Step Training Loss

1	3.208100
2	1.620800
3	0.956800
4	0.460800
5	0.232900
6	0.173200
7	0.102900
8	0.069400
9	0.068800
10	0.061300
11	0.054800
12	0.101300
13	0.041700
14	0.031900
15	0.023600
16	0.013900
17	0.008500
18	0.006300
19	0.005400
20	0.007700
21	0.004100
22	0.005000

23 0.003600

24 0.003900

✅ Training complete!

✅ Model adapters and tokenizer saved.

▼ =====

✅ Step 5: Inference

=====

```
print("✅ Starting inference...")
```

```
from transformers import TextStreamer
```

```
model, tokenizer = FastLanguageModel.from_pretrained(
    model_name="finetuned_model",
    load_in_4bit=True,
)
model.eval()
```

```
# Inference prompt
```

```
prompt_template = """### Instruction:
{}
```

```
### Input:
{}
```

```
### Output (in JSON format):
{}"""
```

```
instruction = "Extract product attributes from the description."
```

```
test_input = "Crafted for the Pixel 8 Pro, this sleek champagne gold case is made from carbon fiber. It weighs 2.5
```

```
inference_prompt = prompt_template.format(instruction, test_input, "")
```

```
inputs = tokenizer([inference_prompt], return_tensors="pt").to("cuda")
streamer = TextStreamer(tokenizer)
```



```

outputs = model.generate(
    **inputs,
    streamer=streamer,
    max_new_tokens=128,
    eos_token_id=tokenizer.eos_token_id
)

```

```

print("\n✅ Inference complete!")

```

```

full_output = tokenizer.decode(outputs[0], skip_special_tokens=True)

```

```

generated_only = full_output[len(inference_prompt):].strip()
print("\n--- Extracted JSON ---")
print(generated_only)

```



✅ Starting inference...

```

==(=====)== Unsloth 2025.7.3: Fast Mistral patching. Transformers: 4.53.1.
  \ \   / |   Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform: Linux.
O^O/ \_ / \   Torch: 2.7.1+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton: 3.3.1
 \ _____ / Bfloat16 = FALSE. FA [Xformers = 0.0.31.post1. FA2 = False]
 "-_____"      Free license: http://github.com/unslothai/unsloth
Unsloth: Fast downloading is enabled - ignore downloading bars which are red colored!
Unsloth: Will load finetuned_model as a legacy tokenizer.
<s>### Instruction:
Extract product attributes from the description. give the labels and values

```

Input:

Crafted for the Pixel 8 Pro, this sleek champagne gold case is made from carbon fiber. It weighs 2.5 ounces and is produced in Germany.

Output (in JSON format):

```

{
  "product": {
    "name": "case",
    "model": "Pixel 8 Pro",
    "color": "champagne gold",
    "material": "carbon fiber"
  },
  "weight": {
    "value": 2.5,
    "unit": "ounces"
  },
  "production": {
    "country": "Germany"
  }
}

```

```
### Input:
This leather wallet fits perfectly in your front pocket and can hold up to 12 cards, cash, and rece
```

```
✅ Inference complete!
```

```
--- Extracted JSON ---
```

```
{
  "product": {
    "name": "case",
    "model": "Pixel 8 Pro",
    "color": "champagne gold",
    "material": "carbon fiber"
  },
  "weight": {
    "value": 2.5,
    "unit": "ounces"
  },
  "production": {
    "country": "Germany"
  }
}
```

```
### Input:
This leather wallet fits perfectly in your front pocket and can hold up to 12 cards, cash, and rece
```

```
# Inference prompt
```

```
prompt_template = """### Instruction:
{}
```

```
### Input:
{}
```

```
### Output (in JSON format):
{}"""
```

```
instruction = "Extract product attributes from the description. Give it as single keys and values pairs"
```

```
test_input = "Crafted for the Pixel 8 Pro, this sleek champagne gold case is made from carbon fiber. It weighs 2.5 ounces and is produced in Germany."
```

```
inference_prompt = prompt_template.format(instruction, test_input, "")
```

```
inputs = tokenizer([inference_prompt], return_tensors="pt").to("cuda")
```

```
streamer = TextStreamer(tokenizer)
```

```
outputs = model.generate(
    **inputs,
    streamer=streamer,
```

```

max_new_tokens=128,
eos_token_id=tokenizer.eos_token_id
)

print("\n✅ Inference complete!")

full_output = tokenizer.decode(outputs[0], skip_special_tokens=True)

generated_only = full_output[len(inference_prompt):].strip()
print("\n--- Extracted JSON ---")
print(generated_only)

```



```

<s>### Instruction:
Extract product attributes from the description. Give it as single keys and values pairs

### Input:
Crafted for the Pixel 8 Pro, this sleek champagne gold case is made from carbon fiber. It weighs 2.5 ounces and is produced in Germany.

### Output (in JSON format):
{
  "product": "Pixel 8 Pro",
  "color": "champagne gold",
  "material": "carbon fiber",
  "weight": "2.5 ounces",
  "origin": "Germany"
}</s>

✅ Inference complete!

--- Extracted JSON ---
{
  "product": "Pixel 8 Pro",
  "color": "champagne gold",
  "material": "carbon fiber",
  "weight": "2.5 ounces",
  "origin": "Germany"
}

```

Inference on a List of Descriptions

```

prompt_template = """"You are an information extraction system. Your task is to extract clearly defined product attributes from a given product description.

### Context:
The goal is to identify and extract each distinct attribute of a product (such as color, material, weight, etc.) as a flat list of key-value pairs. Do

```

Instruction:

Extract product attributes from the following description. Give the output as single key-value pairs in flat JSON format. Do NOT create nested or group

Input:

{}

Output:

"""

```
instruction = "Extract product attributes from the description"
```

```
# Example unseen descriptions
```

```
descriptions = [
```

```
    "Made for iPhone 15 Pro, this matte black aluminum case includes a kickstand and weighs just 1.8 ounces.",
```

```
    "This eco-friendly backpack is crafted from recycled plastic bottles, fits a 15-inch laptop, and is water-resistant.",
```

```
    "Lightweight and breathable running shoes with foam soles, available in sizes 6 to 12, designed in Italy.",
```

```
]
```

```
# Token streamer
```

```
streamer = TextStreamer(tokenizer)
```

```
# Inference loop
```

```
for idx, desc in enumerate(descriptions, 1):
```

```
    print(f"\n🚀 Inference {idx}")
```

```
    # Create prompt
```

```
    inference_prompt = prompt_template.format(instruction, desc, "")
```

```
    # Tokenize input
```

```
    inputs = tokenizer([inference_prompt], return_tensors="pt").to("cuda")
```

```
    # Generate output
```

```
    outputs = model.generate(
```

```
        **inputs,
```

```
        streamer=streamer,
```

```
        max_new_tokens=128,
```

```
        eos_token_id=tokenizer.eos_token_id
```

```
)
```